


LECTURE 17

Random Variables

Numerical functions of random samples and their properties; sampling variability.

The Bias-Variance Tradeoff

What is the mathematical underpinning of this plot?



We'll come back to this...

Our Goal Today

Formalize the notions of **sample statistic, population parameter**.


1. **sample mean** - the mean of your random sample.

`np.mean(data)`

2. The Central Limit Theorem: If you draw a large random sample with replacement, then, regardless of the population distribution, the probability distribution of the sample mean:

- Is roughly normal
- Is centered at the **population mean**
- Has an SD = $\frac{\text{population SD}}{\sqrt{\text{sample size}}}$

We will go over **just enough probability** to help you understand its implications for modeling.



Today's Roadmap

Random Variables and Distributions

Expectation and Variance

Sums of Random Variables

- Equality vs Identically Distributed vs. IID
- Properties of Expectation and Variance
- Covariance, Correlation

Bernoulli and Binomial Random Variables

Sample Statistics

- Sample Mean
- Central Limit Theorem

Derivations

[Terminology] Random Variable

Suppose we draw a random sample of size n from a population.

A **random variable** is a numerical function of a sample.

sample was drawn at random

value depends on how the sample came out

- Often denoted with uppercase “variable-like” letters (e.g. X, Y).
- Also known as a sample statistic, or **statistic**.
- Domain (input): all random samples of size n
- Range (output): number line


[Terminology] Random Variable

Suppose we draw a random sample of size n from a population.

A **random variable** is a **numerical function** of a sample.

- sample was drawn at random
- value depends on how the sample came out
- Often denoted with uppercase “variable-like” letters (e.g. X, Y).
- Also known as a sample statistic, or **statistic**.
- Domain (input): all random samples of size n
- Range (output): number line

Suppose you draw a random sample s of size 3 from the following population:



Define $X = \# \text{ of blue people}$.

X is a random variable!

From Population to Distribution

$X(s)$	
0	3
1	4
2	4
3	6
4	8
...	...
79995	6
79996	6
79997	4
79998	6
...	...

$X(s)$ from all possible samples

$$P(X = x) = \frac{\# \text{ times where } X = x}{\text{pop. size}}$$



x	$P(X = x)$
3	0.1
4	0.2
6	0.4
8	0.3

Probability Distribution Table

[Terminology] Distribution

The **distribution** of a random variable X is a description of how the total probability of 100% is split over all the possible values of X .

A distribution fully defines a random variable.

Assuming (for now) that X is discrete, i.e., has a finite number possible values:

$$P(X = x)$$

The probability that random variable X takes on the value x .


$$\sum_{\text{all } x} P(X = x) = 1 \quad \text{Probabilities must sum to 1.}$$

x	$P(X = x)$
3	0.1
4	0.2
6	0.4
8	0.3




Probabilities are Areas of Histograms

Distribution of **discrete** random variable X



The area of the red bars is
 $P(7 \leq X \leq 9)$.

Distribution of **continuous** random variable Y



The red area under the curve is
 $P(6.8 \leq Y \leq 9.5)$.

Understanding Random Variables

Compute the following probabilities for the random variable X.

1. $P(X = 4) = 0.2$


2. $P(X < 6) = 0.1 + 0.2 = 0.3$

3. $P(X \leq 6) = 0.1 + 0.2 + 0.4 = 0.7$

4. $P(X = 7) = 0$

5. $P(X \leq 8) = 1$

x	$P(X = x)$
3	0.1
4	0.2
6	0.4
8	0.3



Common Random Variables

Bernoulli(p)

- Takes on value 1 with probability p , and 0 with probability $1 - p$
- AKA the “indicator” random variable.



We'll go over these in detail.

The rest are provided for your reference.

Binomial(n, p)

- Number of 1s in n independent Bernoulli(p) trials
- Probabilities given by the binomial formula (Lecture 2)

Uniform on a finite set of values

- Probability of each value is $1 / (\text{size of set})$
- For example, a standard die

Uniform on the unit interval(0, 1)

- Density is flat on (0, 1) and 0 elsewhere

Normal(μ, σ^2)

The numbers in parentheses are the **parameters** of a random variable, which are constants. Parameters define a random variable's shape (i.e., distribution) and its values.

From Distribution to (Simulated) Population

?

$$P(X = x) = \frac{\# \text{ times where } X = x}{\text{pop. size}}$$



x	P(X = x)
3	0.1
4	0.2
6	0.4
8	0.3

Given a random variable's distribution, how could we **generate/simulate** a population?


Probability Distribution Table

$X(s)$ from all possible samples?

From Distribution to (Simulated) Population

X(s)	
0	3
1	4
2	4
3	6
4	8
...	...
79995	6
79996	6
79997	4
79998	6
...	...

$$P(X = x) = \frac{\# \text{ times where } X = x}{\text{pop. size}}$$



x	P(X = x)
3	0.1
4	0.2
6	0.4
8	0.3

Probability
Distribution Table

Simulate: Randomly pick values of X according to its distribution

`np.random.choice` or `df.sample`

X(s) from many, many
(simulated) samples

Expectation and Variance

Random Variables and Distributions

Expectation and Variance

Sums of Random Variables

- Equality vs Identically Distributed vs. IID
- Properties of Expectation and Variance
- Covariance, Correlation


Bernoulli and Binomial Random Variables

Sample Statistics

- Sample Mean
- Central Limit Theorem

Descriptive Properties of Random Variables

There are several ways to describe a random variable:



Definition of Expectation

The **expectation** of a random variable X is the **weighted average** of the values of X , where the weights are the probabilities of the values.

Two equivalent ways to apply the weights:

1. One sample at a time:

$$\mathbb{E}[X] = \sum_{\text{all samples } s} X(s)P(s)$$

1. One possible value at a time:

$$\mathbb{E}[X] = \sum_{\text{all possible } x} xP(X = x)$$

} More common (we are usually given the distribution, not all possible samples)

Expectation is a **number, not a **random variable**!**

- It is analogous to the **average** (same units as the random variable).
- It is the center of gravity of the probability histogram.
- It is the long run average of the random variable, if you simulate the variable many times.

Example

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

Consider the random variable X we defined earlier.

$$\mathbb{E}[X] = \sum_x x \cdot P(X = x)$$

$$= 3 \cdot 0.1 + 4 \cdot 0.2 + 6 \cdot 0.4 + 8 \cdot 0.3$$

$$= 0.3 + 0.8 + 2.4 + 2.4$$


$$= 5.9$$

Note, $E[X] = 5.9$ is not a possible value of X !

It is an average.

The expectation of X does not need to be a value of X .

x	P(X = x)
3	0.1
4	0.2
6	0.4
8	0.3



Definition of Variance

Variance is the **expected squared deviation from the expectation** of X.

$$\text{Var}(X) = \mathbb{E} [(X - \mathbb{E}[X])^2]$$

- The units of the variance are the square of the units of X.
- To get back to the right scale, use the **standard deviation** of X: $\text{SD}(X) = \sqrt{\text{Var}(X)}$

Variance is a number, not a random variable!

- The main use of variance is to **quantify chance error**. How far away from the expectation could X be, just by chance?

By [Chebyshev's inequality](#) :

- No matter what the shape of the distribution of X is, the vast majority of the probability lies in the interval "expectation plus or minus a few SDs."

There's a more convenient form of variance:

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$$

- Proof (involves expanding the square and properties of expectation/summations): [link](#)
- Useful in Mean Squared Error calculations
 - If X is centered (i.e. $\mathbb{E}[X] = 0$), then $\mathbb{E}[X^2] = \text{Var}(X)$
- When computing variance by hand, often used instead of definition.

Rolling a Die

Let X be the outcome of a single die roll.
 X is a random variable.

$$P(X = x) = \begin{cases} 1/6 & \text{if } x \in \{1, 2, 3, 4, 5, 6\} \\ 0 & \text{otherwise} \end{cases}$$



1. What is the expectation, $E[X]$?

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\begin{aligned}\text{Var}(X) &= \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2\end{aligned}$$

(definitions/properties)

2. What is the variance,
 $\text{Var}(X)$?



Dice Is the Plural; Die Is the Singular

Let X be the outcome of a single die roll.
 X is a random variable.

$$P(X = x) = \begin{cases} 1/6 & \text{if } x \in \{1, 2, 3, 4, 5, 6\} \\ 0 & \text{otherwise} \end{cases}$$



1. What is the expectation, $E[X]$?

$$\begin{aligned}\mathbb{E}[X] &= 1(1/6) + 2(1/6) + 3(1/6) + 4(1/6) + 5(1/6) + 6(1/6) \\ &= (1/6)(1 + 2 + 3 + 4 + 5 + 6) = \frac{7}{2}\end{aligned}$$

2. What is the variance,
 $\text{Var}(X)$?

$$\begin{aligned}\mathbb{E}[X] &= \sum_x x P(X = x) \\ \text{Var}(X) &= \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2\end{aligned}$$

Dice Is the Plural; Die Is the Singular

Let X be the outcome of a single die roll.
 X is a random variable.

$$P(X = x) = \begin{cases} 1/6 & \text{if } x \in \{1, 2, 3, 4, 5, 6\} \\ 0 & \text{otherwise} \end{cases}$$



1. What is the expectation, $E[X]$?

$$\begin{aligned} \mathbb{E}[X] &= 1(1/6) + 2(1/6) + 3(1/6) + 4(1/6) + 5(1/6) + 6(1/6) \\ &= (1/6)(1 + 2 + 3 + 4 + 5 + 6) = \frac{7}{2} \end{aligned}$$

$$\begin{aligned} \mathbb{E}[X] &= \sum_x x P(X = x) \\ \text{Var}(X) &= \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2 \end{aligned}$$

2. What is the variance,
 $\text{Var}(X)$?

Approach 1: Definition

$$\begin{aligned} \text{Var}(X) &= (1/6) ((1 - 7/2)^2 + (2 - 7/2)^2 \\ &\quad + (3 - 7/2)^2 + (4 - 7/2)^2 \\ &\quad + (5 - 7/2)^2 + (6 - 7/2)^2) \\ &= 35/12 \end{aligned}$$

Approach 2: Property

$$\mathbb{E}[X^2] = \sum_x x^2 P(X = x)$$

$$1^2 * (1/6) + 2^2 * (1/6) + \dots + 6^2 * (1/6) = 91/6$$

$$\text{Var}(X) = 91/6 - (7/2)^2 = 35/12$$

Sums of Random Variables

Random Variables and Distributions
Expectation and Variance

Sums of Random Variables

- Equality vs Identically Distributed vs. IID
- Properties of Expectation and Variance
- Covariance, Correlation

Bernoulli and Binomial Random Variables
Sample Statistics

- Sample Mean
- Central Limit Theorem

Functions of Multiple Random Variables

A function of a random variable is also a random variable!

If you create multiple random variables based on your sample...

...then functions of those random variables are also random variables.

For instance, if X_1, X_2, \dots, X_n are random variables, then so are all of these:

$$X_n^2 \quad \#\{i : X_i > 10\}$$

$$\max(X_1, X_2, \dots, X_n) \quad \frac{1}{n} \sum_{i=1}^n (X_i - c)^2$$

$$\frac{1}{n} \sum_{i=1}^n X_i$$

Many functions of RVs that we care about (**counts, means**) involve **sums of RVs**, so we expand on properties of sums of RVs.

Equal vs. Identically Distributed vs. IID

Suppose that we have two random variables X and Y .

X and Y are **equal** if:


- $X(s) = Y(s)$ for every sample s .
- We write $X = Y$.

X and Y are **identically distributed** if:

- The distribution of X is the same as the distribution of Y
- We say " X and Y are equal in distribution."
- If $X = Y$, then X and Y are identically distributed; but the converse is not true.

X and Y are **independent and identically distributed (IID)** if:

- X and Y are identically distributed, and
- Knowing the outcome of X does not influence your belief of the outcome of Y , and vice versa (" X and Y are independent.")



Equal RVs

IID RVs

Distributions of Sums



Demo

Let X_1 and X_2 be numbers on two rolls of a die.



- X_1, X_2 are **IID**, so X_1, X_2 have the same distribution.
- But the sums $\mathbf{Y} = \mathbf{X}_1 + \mathbf{X}_2 = 2\mathbf{X}_1$ and $\mathbf{Z} = \mathbf{X}_1 + \mathbf{X}_2$ have different distributions!

Let's show this through simulation:



- Same expectation...
- But $\mathbf{Y} = 2\mathbf{X}_1$ has larger variance!

How can we directly compute $E[Y]$, $\text{Var}(Y)$, **without** simulating distributions?

$E[\cdot]$	6.984400	6.984950
$\text{Var}(\cdot)$	11.669203	5.817246
$\text{SD}(\cdot)$	3.416021	2.411897

Properties of Expectation [1/3]

Instead of simulating full distributions, we often just compute expectation and variance directly.

Recall definitions of expectation:

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\mathbb{E}[X] = \sum_{\text{all samples}} X(s) P(s)$$

Properties:

1. **Expectation is linear.**

Intuition: summations are linear. [Proof](#)

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

Properties of Expectation [2/3]

Instead of simulating full distributions, we often just compute expectation and variance directly.

Recall definitions of expectation:

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\mathbb{E}[X] = \sum_{\text{all samples}} X(s) P(s)$$

Properties:

1. **Expectation is linear.**

Intuition: summations are linear. [Proof](#)

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

2. Expectation is linear in sums of RVs,
for any relationship between X and Y. [Proof](#)

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

Properties of Expectation [3/3]

Instead of simulating full distributions, we often just compute expectation and variance directly.

Recall definitions of expectation:

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\mathbb{E}[X] = \sum_{\text{all samples}} X(s) P(s)$$

Properties:

1. **Expectation is linear.**

Intuition: summations are linear. [Proof](#)

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

2. Expectation is linear in sums of RVs,
for any relationship between X and Y. [Proof](#)

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

3. If g is a non-linear function, then in general $\mathbb{E}[g(X)] \neq g(\mathbb{E}[X])$.

- Example: if X is -1 or 1 with equal probability, then $E[X] = 0$ but $E[X^2] = 1 \neq 0$.

Properties of Variance [1/2]

Recall definition of variance:

$$\text{Var}(X) = \mathbb{E} [(X - \mathbb{E}[X])^2]$$




Properties:

1. Variance is non-linear:

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

Intuition ([full proof](#)): Consider the Standard Deviation for $Y = -3X + 2$:

$$\text{SD}(aX + b) = |a| \text{SD}(X)$$



Properties of Variance [2/2]

Recall definition of variance:

$$\text{Var}(X) = \mathbb{E} [(X - \mathbb{E}[X])^2]$$

Properties:

1. Variance is non-linear:

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

Intuition (full proof): Consider the Standard Deviation for $Y = -3X + 2$:

$$\text{SD}(aX + b) = |a| \text{SD}(X)$$

2. Variance of sums of RVs is affected by the (in)dependence of the RVs ([derivation](#)):

$$\text{Var}(X_1 + X_2) = \text{Var}(X_1) + \text{Var}(X_2) + 2\text{Cov}(X, Y)$$



Covariance of X and Y (next slide).
If X, Y independent,
then $\text{Cov}(X, Y) = 0$.

Covariance and Correlation: The Basics

Covariance is the expected product of deviations from expectation.

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

- A generalization of variance. Note $\text{Cov}(X, X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \text{Var}(X)$.
- Interpret by defining **correlation** (yes, *that* correlation!):

$$r(X, Y) = \mathbb{E}\left[\underbrace{\left(\frac{X - \mathbb{E}[X]}{\text{SD}(X)}\right)}_{\text{standard units of } X \text{ (link)}} \left(\frac{Y - \mathbb{E}[Y]}{\text{SD}(Y)}\right)\right] = \frac{\text{Cov}(X, Y)}{\text{SD}(X)\text{SD}(Y)}$$

Correlation (and therefore covariance) measures a linear relationship between X and Y.

Covariance and Correlation: The Basics

Covariance is the expected product of deviations from expectation.

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

- A generalization of variance. Note $\text{Cov}(X, X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \text{Var}(X)$.
- Interpret by defining **correlation** (yes, *that* correlation!):

$$r(X, Y) = \mathbb{E}\left[\left(\frac{X - \mathbb{E}[X]}{\text{SD}(X)}\right)\left(\frac{Y - \mathbb{E}[Y]}{\text{SD}(Y)}\right)\right] = \frac{\text{Cov}(X, Y)}{\text{SD}(X)\text{SD}(Y)}$$

standard units of X ([link](#))

Correlation (and therefore covariance) measures a linear relationship between X and Y.

- If X and Y are correlated, then knowing X tells you something about Y.
- “X and Y are uncorrelated” is the same as “Correlation and covariance equal to 0”
- **Independent X, Y are uncorrelated**, because knowing X tells you nothing about Y.
- The converse is not necessarily true: **X, Y could be uncorrelated but not independent**.


Dice, Our Old Friends: Expectation

Let X_1 and X_2 be numbers on two rolls of a die.


- X_1, X_2 are **IID**, so X_1, X_2 have the same distribution.
- Therefore $E[X_1] = E[X_2] = 7/2$ $\text{Var}(X_1) = \text{Var}(X_2) = 35/12$



$$Y = 2X_1$$



$$Z = X_1 + X_2$$



$$E[Y] = E[2X_1] = 2E[X_1] = 7$$

$$E[Z] = E[X_1] + E[X_2] = (7/2) + (7/2) = 7$$


Dice, Our Old Friends: Variance

Let X_1 and X_2 be numbers on two rolls of a die.

- X_1, X_2 are **IID**, so X_1, X_2 have the same distribution.
- Therefore $E[X_1] = E[X_2] = 7/2$ $\text{Var}(X_1) = \text{Var}(X_2) = 35/12$




$$Y = 2X_1$$



$$E[Y] = E[2X_1] = 2E[X_1] = 7$$

$$\begin{aligned}\text{Var}(Y) &= \text{Var}(2X_1) = 4\text{Var}(X_1) \\ &= 4(35/12) \\ &\approx 11.67\end{aligned}$$

$$Z = X_1 + X_2$$



$$\begin{aligned}E[Z] &= E[X_1] + E[X_2] = (7/2) + (7/2) = \\ &7\end{aligned}$$

$$\begin{aligned}\text{Var}(Z) &= \text{Var}(X_1) + \text{Var}(X_2) + 2\text{Cov}(X_1, X_2) \\ &= (35/12) + (35/12) + 0 \\ &\approx 5.83\end{aligned}$$

0
 X_1, X_2 independent

[Summary] Expectation and Variance for Linear Functions of Random Variables



Let X be
a random variable with
distribution $P(X = x)$.

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\text{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] \quad (\text{definition})$$

$$= \mathbb{E}[X^2] - (\mathbb{E}[X])^2 \quad (\text{easier computation})$$

Let a and b be
scalar values.

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

Let Y be
another random variable.

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

$$\text{Var}(X_1 + X_2) = \text{Var}(X_1) + \text{Var}(X_2) + \underbrace{2\text{Cov}(X, Y)}$$

Zero if X, Y independent.

Bernoulli and Binomial Random Variables

Random Variables and Distributions

Expectation and Variance

Sums of Random Variables

- Equality vs Identically Distributed vs. IID
- Properties of Expectation and Variance
- Covariance, Correlation

Bernoulli and Binomial Random Variables

Sample Statistics

- Sample Mean
- Central Limit Theorem

Common Random Variables

Bernoulli(p)

- Takes on value 1 with probability p , and 0 with probability $1 - p$
- AKA the “indicator” random variable.

Binomial(n, p)

- Number of 1s in n independent Bernoulli(p) trials
- Probabilities given by the binomial formula (Lecture 2)

Uniform on a finite set of values

- Probability of each value is $1 / (\text{size of set})$
- For example, a standard die

Uniform on the unit interval(0, 1)

- Density is flat on $(0, 1)$ and 0 elsewhere

Normal(μ, σ^2)



We'll now revisit these to solidify our understanding of expectation/variance.

Properties of Bernoulli Random Variables

Let X be a **Bernoulli**(p) random variable.

- Takes on value 1 with probability p , and 0 with probability $1 - p$
- AKA the “indicator” random variable.

$$\mathbb{E}[X] = \sum_x xP(X = x)$$

$$\begin{aligned}\text{Var}(X) &= \mathbb{E}[(X - \mathbb{E}[X])^2] \\ &= \mathbb{E}[X^2] - (\mathbb{E}[X])^2\end{aligned}$$

Definitions

Expectation:

$$\mathbb{E}[X] = 1 \cdot p + 0 \cdot (1 - p) = p$$

We will get an average value of p across many, many samples

Variance:

$$\mathbb{E}[X^2] = 1^2 \cdot p + 0 \cdot (1 - p) = p$$

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$$

$$= p - p^2 = p(1 - p)$$

Lower Var: $p = 0.1$ or 0.9
Higher Var: p close to 0.5

Properties of Binomial Random Variables

Let Y be a **Binomial**(n , p) random variable.

- Y is the number (i.e., count) of 1s in n independent Bernoulli(p) trials.
- Distribution of Y given by the binomial formula (Lecture 2).

We can write:
$$Y = \sum_{i=1}^n X_i$$

A count is a **sum** of 0's and 1's.

- X_i is the indicator of success on trial i . $X_i = 1$ if trial i is a success, else 0.
- All X_i 's are **IID** (independent and identically distributed) and **Bernoulli**(p).

Properties of Binomial Random Variables

Let Y be a **Binomial**(n, p) random variable.

- Y is the number (i.e., count) of 1s in n independent Bernoulli(p) trials.
- Distribution of Y given by the binomial formula (Lecture 2).

We can write:
$$Y = \sum_{i=1}^n X_i$$

A count is a sum of 0's and 1's.

- X_i is the indicator of success on trial i . $X_i = 1$ if trial i is a success, else 0.
- All X_i 's are **IID** (independent and identically distributed) and **Bernoulli**(p).

Expectation:
$$\mathbb{E}[Y] = \sum_{i=1}^n \mathbb{E}[X_i] = np$$

Variance: Because all X_i 's are independent, $\text{Cov}(X_i, X_j) = 0$ for all i, j .

$$\text{Var}(Y) = \sum_{i=1}^n \text{Var}(X_i) = np(1 - p)$$

Interlude

Suppose you win cash based on the number of heads you get in a series of 20 coin flips.

Let $X_i = 1$ if the i -th coin is heads, 0 otherwise

Which payout strategy would you choose?

Hint: Compare expectations and variances.

A. $Y_A = 10 \cdot X_1 + 10 \cdot X_2$

B. $Y_B = \left(\sum_{i=1}^{20} X_i \right)$

C. $Y_C = 20 \cdot X_1$

Which would you pick?

Suppose you win cash based on the number of heads you get in a series of 20 coin flips.

Let X_1, X_2, \dots, X_{20} be 20 **IID** Bernoulli(0.5) random variables.

- Since X_i s are independent: $\text{Cov}(X_i, X_j) = 0$ for all i, j.
- Since X_i is Bernoulli($p = 0.5$): $E[X_i] = p = 0.5$, $\text{Var}(X_i) = p(1-p) = 0.25$.

Expectation of a linear function is linear

$$\text{Var}(aX + b) = a^2\text{Var}(X)$$

$$\text{Var}(X_1 + X_2) = \text{Var}(X_1) + \text{Var}(X_2) + 2\text{Cov}(X, Y)$$

Which payout strategy would you choose?

	A. $Y_A = 10 \cdot X_1 + 10 \cdot X_2$	B. $Y_B = \left(\sum_{i=1}^{20} X_i \right)$	C. $Y_C = 20 \cdot X_1$
Expectation			
Variance			
Std. Deviation			



Which would you pick?

Suppose you win cash based on the number of heads you get in a series of 20 coin flips.

Let X_1, X_2, \dots, X_{20} be 20 **IID** Bernoulli(0.5) random variables.

- Since X_i s are independent: $\text{Cov}(X_i, X_j) = 0$ for all i, j.
- Since X_i is Bernoulli($p = 0.5$): $E[X_i] = p = 0.5$, $\text{Var}(X_i) = p(1-p) = 0.25$.

Expectation of a linear function is linear

$$\text{Var}(aX + b) = a^2\text{Var}(X)$$


$$\text{Var}(X_1 + X_2) = \text{Var}(X_1) + \text{Var}(X_2) + 2\text{Cov}(X, Y)$$

Which payout strategy would you choose?

	A. $Y_A = \$10 \cdot X_1 + \$10 \cdot X_2$	B. $Y_B = \$\left(\sum_{i=1}^{20} X_i\right)$	C. $Y_C = \$20 \cdot X_1$
Expectation	$E[Y_A] = 10(0.5) + 10(0.5) = 10$	$E[Y_B] = 0.5 + \dots + 0.5 = 10$	$E[Y_C] = 20(0.5) = 10$
Variance	$\text{Var}(Y_A) = 10^2(0.25) + 10^2(0.25) = 50$	$\text{Var}(Y_B) = 0.25 + \dots + 0.25 = 20(0.25) = 5$	$\text{Var}(Y_C) = 20^2(0.25) = 100$
Std. Deviation	$\text{SD}(Y_A) \approx 7.07$	$\text{SD}(Y_B) \approx 2.24$	$\text{SD}(Y_C) = 10$

Binomial(n , p) for large n

For $p = 0.5$, $n = 50$ (i.e. number of heads in 50 fair coin flips):



Sample Statistics

Random Variables and Distributions

Expectation and Variance

Sums of Random Variables

- Equality vs Identically Distributed vs. IID
- Properties of Expectation and Variance
- Covariance, Correlation

Bernoulli and Binomial Random Variables

Sample Statistics

- Sample Mean
- Central Limit Theorem

From Populations to Samples


Today, we've talked extensively about **populations**:

- If we know the **distribution of a random variable**, we can reliably compute expectation, variance, functions of the random variable, etc.


However, in Data Science, we often collect **samples**.

- We don't know the distribution of our population.
- We'd like to use the distribution of your sample to estimate/infer properties of the population.

The **big assumption** we make in modeling/inference:



The Sample is a Set of IID Random Variables



Or

Population (really large N)

79995	6
79996	6
79997	4
79998	6
79999	6
...	...

Sample at Random with Replacement




`df.sample(n,
replace=True)`
[\[documentation\]](#)

Each observation in our sample is a **Random Variable** drawn **IID** from our population distribution.

Sample
($n \ll N$)

X_1, X_2, \dots, X_n

The Sample is a Set of IID Random Variables



Sample at Random
with Replacement



x	$P(X = x)$	$X(s)$	
3	0.1	0	3
4	0.2	1	4
6	0.4	2	4
8	0.3	3	6
		4	8
	
		79995	6
		79996	6
		79997	4
		79998	6
		79999	6
	

$$E[X] = 5.9$$

Population Mean

A **number**,
i.e., fixed value

`df.sample(n,
replace=True)
\[documentation\]`

x
0 6
1 8
2 6
3 6
4 3
...
95 8
96 6
97 6
98 3
99 8

Sample Mean
A **random variable**!

Depends on our
randomly drawn sample!!

$$\text{np.mean}(\dots) = 5.71$$

Sample X_1, X_2, \dots, X_n

[Terminology] Sample Mean

Consider an IID sample X_1, X_2, \dots, X_n drawn from a numerical population with mean μ and SD σ .

Define the **sample mean**:

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

Expectation:

$$\begin{aligned}\mathbb{E}[\bar{X}_n] &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] \\ &= \frac{1}{n}(n\mu) = \mu\end{aligned}$$

Variance/Standard Deviation:

$$\begin{aligned}\text{Var}(\bar{X}_n) &= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) = \frac{1}{n^2} \underbrace{\left(\sum_{i=1}^n \text{Var}(X_i)\right)}_{\text{IID} \rightarrow \text{Cov}(X_i, X_j) = 0} \\ &= \frac{1}{n^2} (n\sigma^2) = \frac{\sigma^2}{n}\end{aligned}$$

$$\text{SD}(\bar{X}_n) = \frac{\sigma}{\sqrt{n}}$$


Distribution?

\bar{X}_n is **normally distributed** by the **Central Limit Theorem**.

Central Limit Theorem

No matter what population you are drawing from:

If an IID sample of size n is large,
the probability distribution of the **sample mean**
is **roughly normal** with mean μ and SD σ/\sqrt{n} .



Any theorem that provides the rough distribution of a statistic
and **doesn't need the distribution of the population** is valuable to data scientists.


- Because we rarely know a lot about the population!

How Large Is “Large”?

No matter what population you are drawing from:

If an IID **sample of size n is large**,

the probability distribution of the sample mean
is **roughly normal** with mean μ and $SD \sigma / \sqrt{n}$.




How large does n have to be for the normal approximation to be good?

- ...It depends on the shape of the distribution of the population...
- If population is **roughly symmetric and unimodal**/uniform, could need as few as **$n = 20$** .
If population is very skewed, you will need bigger n .
- If in doubt, you can bootstrap the sample mean and see if the bootstrapped distribution is bell-shaped.

Accuracy and Spread of the Sample Mean

Our goal is often to **estimate** some characteristic of a population.

- Example: average height of Cal undergraduates.
- We typically can collect a **single sample**. It has just one average.
- Since that sample was random, it *could have* come out differently.



We should consider the **average value and spread** of all possible sample means, and what this means for how big n should be.


$$\mathbb{E}[\bar{X}_n] = \mu$$

$$\text{SD}(\bar{X}_n) = \frac{\sigma}{\sqrt{n}}$$

Accuracy and Spread of the Sample Mean

Our goal is often to **estimate** some characteristic of a population.

- Example: average height of Cal undergraduates.
- We typically can collect a **single sample**. It has just one average.
- Since that sample was random, it *could have come out differently*.



We should consider the **average value and spread** of all possible sample means, and what this means for how big n should be.

$$\mathbb{E}[\bar{X}_n] = \mu$$

$$\text{SD}(\bar{X}_n) = \frac{\sigma}{\sqrt{n}}$$

For every sample size, the expected value of the sample mean is the population mean.

We call the sample mean an **unbiased estimator** of the population mean.

Accuracy and Spread of the Sample Mean

Our goal is often to **estimate** some characteristic of a population.


- Example: average height of Cal undergraduates.
- We typically can collect a **single sample**. It has just one average.
- Since that sample was random, it *could have come out differently*.

We should consider the **average value and spread** of all possible sample means, and what this means for how big n should be.

$$\mathbb{E}[\bar{X}_n] = \mu$$

For every sample size, the expected value of the sample mean is the population mean.

We call the sample mean an **unbiased estimator** of the population mean.
(more in next lecture)



$$\text{SD}(\bar{X}_n) = \frac{\sigma}{\sqrt{n}}$$

Square root law: If you increase the sample size by a factor, the SD decreases by the square root of the factor.

The sample mean is more likely to be close to the population mean if we have a larger sample size.

[Extra Slides] Derivations

Random Variables and Distributions

- Expectation and Variance
- Equality vs Identically Distributed
- Common RVs: Bernoulli, Binomial

Functions of Random Variables

- Distributions through Simulation, I.I.D.
- Properties of Expectation and Variance
- Covariance, Correlation
- Standard Units

Sample Statistics

- Sample Mean
- Central Limit Theorem

X in **standard units** is the random variable

$$X_{su} = \frac{X - \mathbb{E}(X)}{\text{SD}(X)}.$$

X_{su} measures X on the scale “**number of SDs from expectation.**”

- It is a linear transformation of X . By the linear transformation rules for expectation and variance:

$$\mathbb{E}(X_{su}) = 0, \quad \text{SD}(X_{su}) = 1$$

- Since X_{su} is centered (has expectation 0):

$$\mathbb{E}(X_{su}^2) = \text{Var}(X_{su}) = 1$$

You should prove these facts yourself.

There's a more convenient form of variance for use in calculations.

$$\text{Var}(X) = \mathbb{E}(X^2) - (\mathbb{E}(X))^2$$

To derive this, we make repeated use of the linearity of expectation.

$$\begin{aligned}\text{Var}(X) &= E((X - E(X))^2) \\&= E(X^2 - 2XE(X) + (E(X))^2) \\&= E(X^2) - 2E(X)E(X) + (E(X))^2 \\&= E(X^2) - (\mathbb{E}(X))^2\end{aligned}$$

Recall definition of expectation:

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\mathbb{E}[X] = \sum_{\text{all samples}} X(s) P(s)$$

1. **Expectation is linear:**

(intuition: summations are linear)

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

Proof:

$$\begin{aligned}\mathbb{E}[aX + b] &= \sum_x (ax + b) P(X = x) = \sum_x (axP(X = x) + bP(X = x)) \\ &= a \sum_x x P(X = x) + b \sum_x P(X = x) \\ &= a\mathbb{E}[X] + b \cdot 1\end{aligned}$$

Recall definitions of expectation:

$$\mathbb{E}[X] = \sum_x x P(X = x)$$

$$\mathbb{E}[X] = \sum_{\text{all samples}} X(s) P(s)$$

3. **Expectation is linear in sums of RVs:**

For any relationship between X and Y.

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b$$

Proof:

$$\begin{aligned}\mathbb{E}[X + Y] &= \sum_s (X + Y)(s) P(s) = \sum_s (X(s) + Y(s)) P(s) \\ &= \sum_s (X(s)P(s) + Y(s)P(s)) \\ &= \sum_s X(s) P(s) + \sum_s Y(s) P(s) \\ &= \mathbb{E}[X] + \mathbb{E}[Y]\end{aligned}$$

We know that $\mathbb{E}(aX + b) = a\mathbb{E}(X) + b$

In order to compute $\text{Var}(aX + b)$, consider:

- A shift by **b** units **does not** affect spread. Thus, $\text{Var}(aX + b) = \text{Var}(aX)$.
- The multiplication by **a** **does** affect spread!

Then,

$$\begin{aligned}\text{Var}(aX + b) &= \text{Var}(aX) = E((aX)^2) - (E(aX))^2 \\ &= E(a^2 X^2) - (aE(X))^2 \\ &= a^2(E(X^2) - (E(X))^2) \\ &= a^2 \text{Var}(X)\end{aligned}$$

In summary:

$$\mathbb{E}[g(X)] \neq g(\mathbb{E}[X])$$

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

$$\text{SD}(aX + b) = |a| \text{SD}(X)$$

Don't forget the absolute values and squares!

The variance of a sum is affected by the dependence between the two random variables that are being added. Let's expand out the definition of $\text{Var}(X + Y)$ to see what's going on.

Let $\mu_x = E[X], \mu_y = E[Y]$

$$\begin{aligned} \text{Var}(X + Y) &= E[(X + Y - E(X + Y))^2] \\ &= E[((X - \mu_x) + (Y - \mu_y))^2] \\ &= E[(X - \mu_x)^2 + 2(X - \mu_x)(Y - \mu_y) + (Y - \mu_y)^2] \\ &= E[(X - \mu_x)^2] + E[(Y - \mu_y)^2] + 2E[(X - \mu_x)(Y - \mu_y)] \\ &= \text{Var}(X) + \text{Var}(Y) + 2E[(X - E(X))(Y - E(Y))] \end{aligned}$$

By the linearity of expectation,
and the substitution.

We see

Addition rule for variance

If X and Y are **uncorrelated** (in particular, if they are **independent**),
then

$$\mathbb{V}ar(X + Y) = \mathbb{V}ar(X) + \mathbb{V}ar(Y)$$

Therefore, under the same conditions,

$$\text{SD}(X + Y) = \sqrt{\mathbb{V}ar(X) + \mathbb{V}ar(Y)} = \sqrt{(\text{SD}(X))^2 + (\text{SD}(Y))^2}$$

- Think of this as “Pythagorean theorem” for random variables.
- Uncorrelated random variables are like orthogonal vectors.